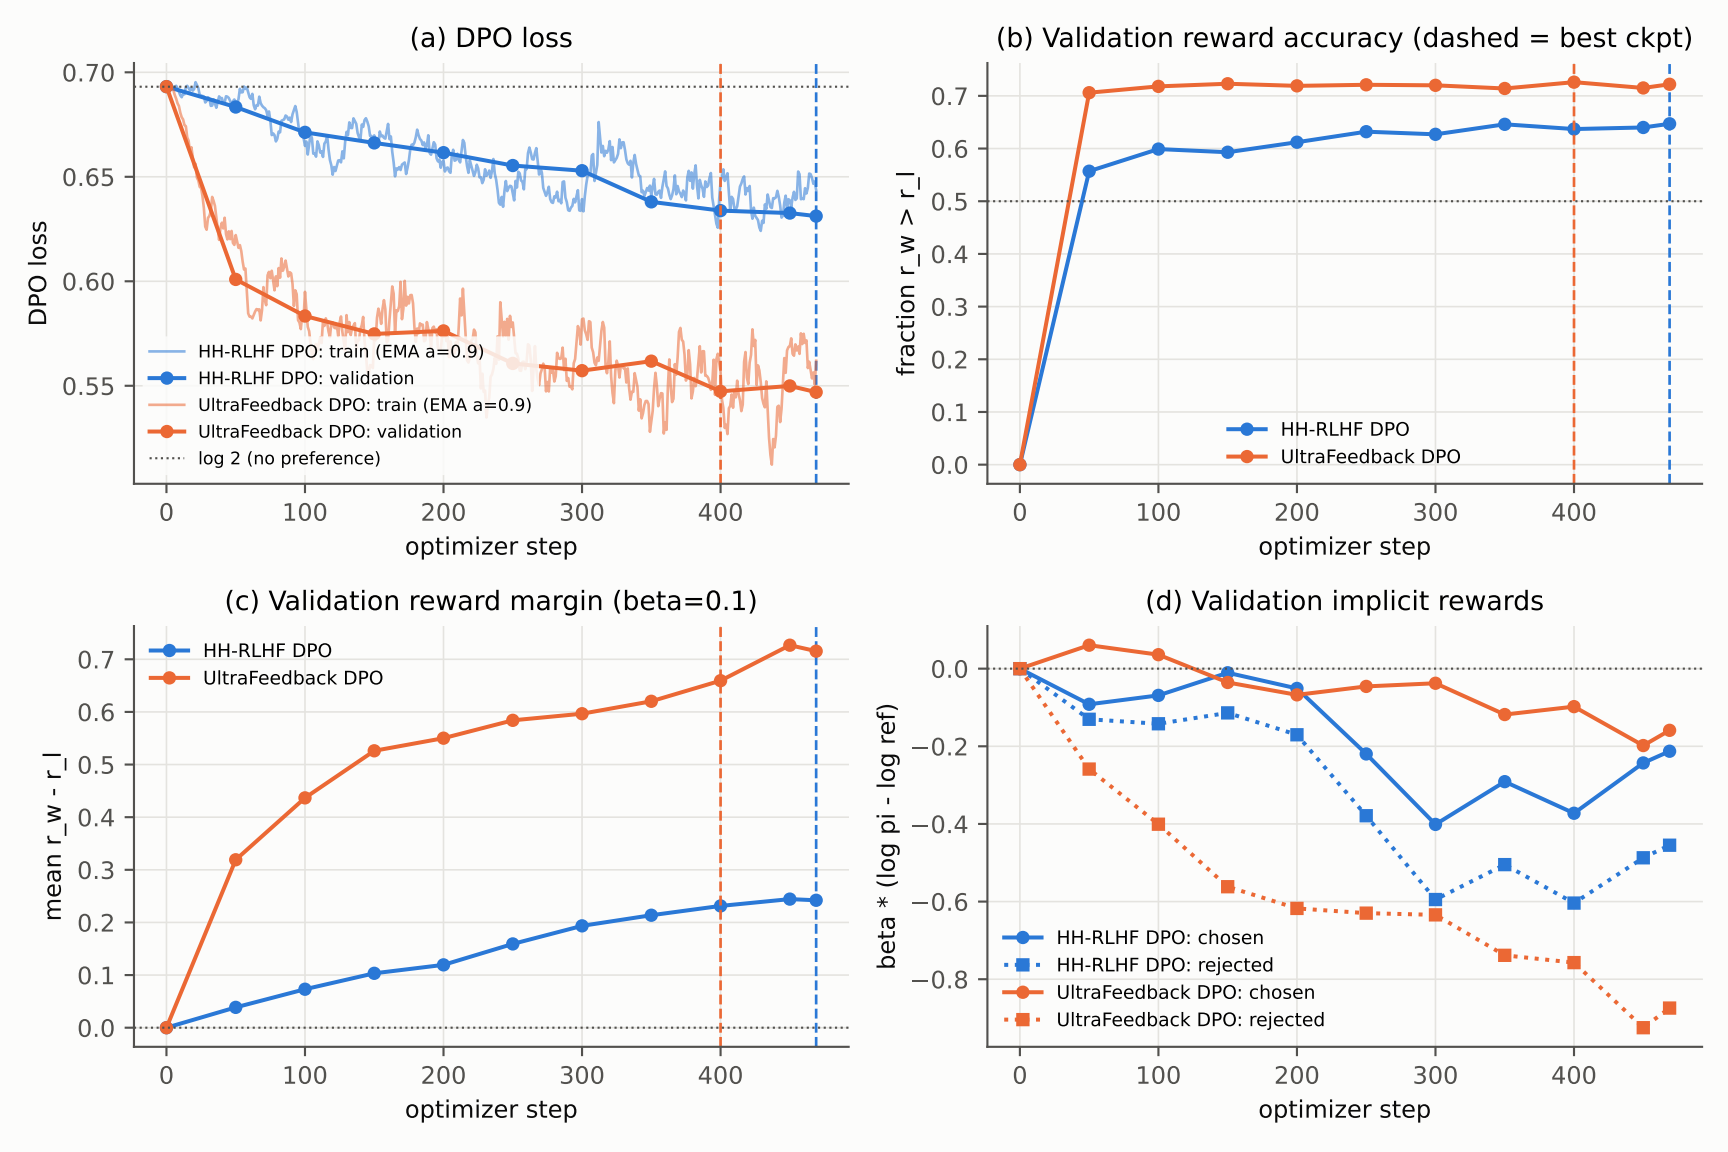
(a) DPO loss
DPO loss
optimizer step
HH-RLHF DPO: train (EMA a=0.9)
HH-RLHF DPO: validation
UltraFeedback DPO: train (EMA a=0.9)
UltraFeedback DPO: validation
log 2 (no preference)
(b) Validation reward accuracy (dashed = best ckpt)
fraction r_w > r_l
optimizer step
HH-RLHF DPO
UltraFeedback DPO
(c) Validation reward margin (beta=0.1)
mean r_w - r_l
optimizer step
HH-RLHF DPO
UltraFeedback DPO
(d) Validation implicit rewards
beta * (log pi - log ref)
optimizer step
HH-RLHF DPO: chosen
HH-RLHF DPO: rejected
UltraFeedback DPO: chosen
UltraFeedback DPO: rejected