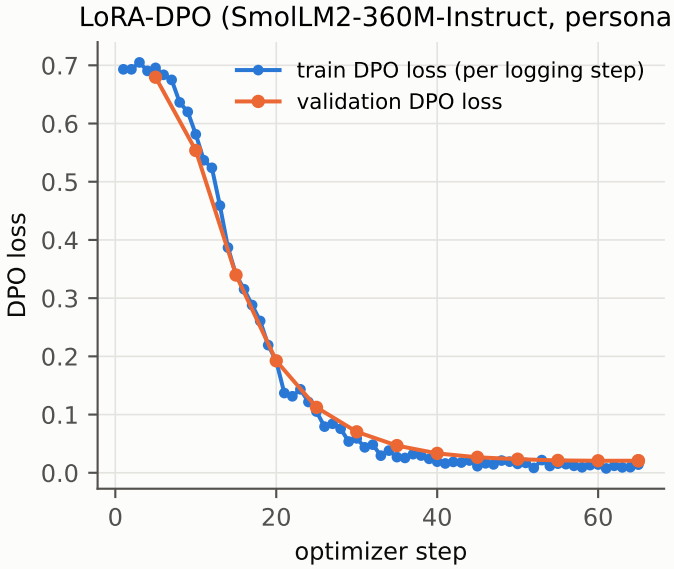
LoRA-DPO (SmolLM2-360M-Instruct, persona)
train DPO loss (per logging step)
validation DPO loss
DPO loss
optimizer step
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0
20
40
60

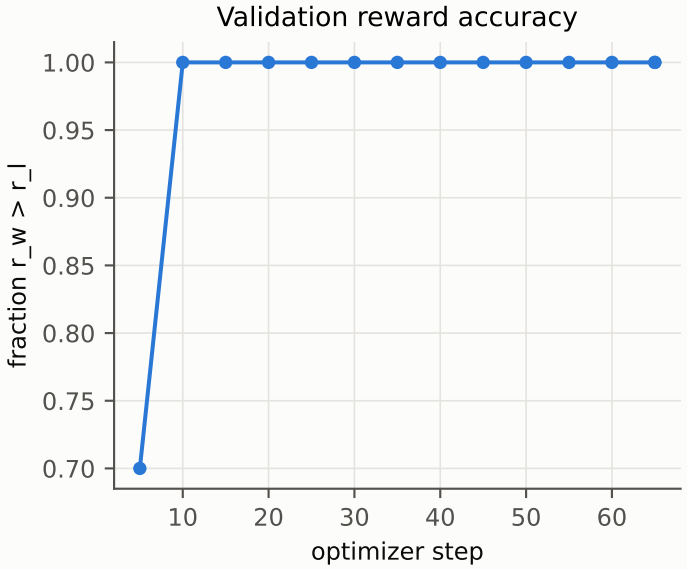
Validation reward accuracy
fraction r_w > r_l
optimizer step
0.70
0.75
0.80
0.85
0.90
0.95
1.00
10
20
30
40
50
60

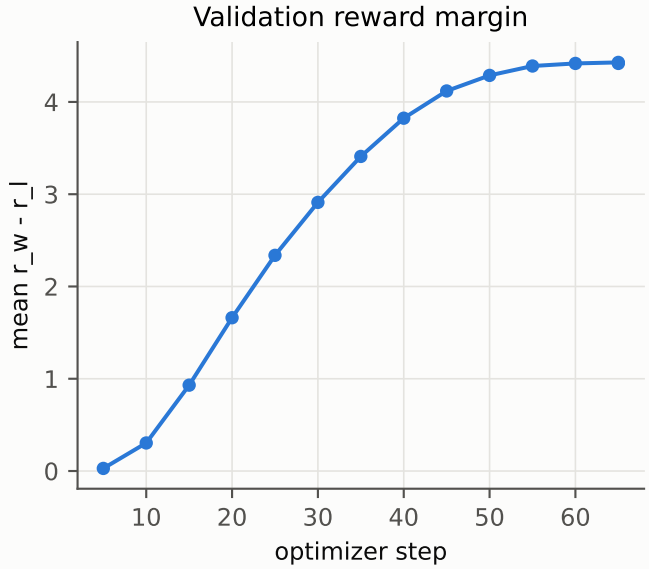
Validation reward margin
mean r_w - r_l
optimizer step
0
1
2
3
4
10
20
30
40
50
60