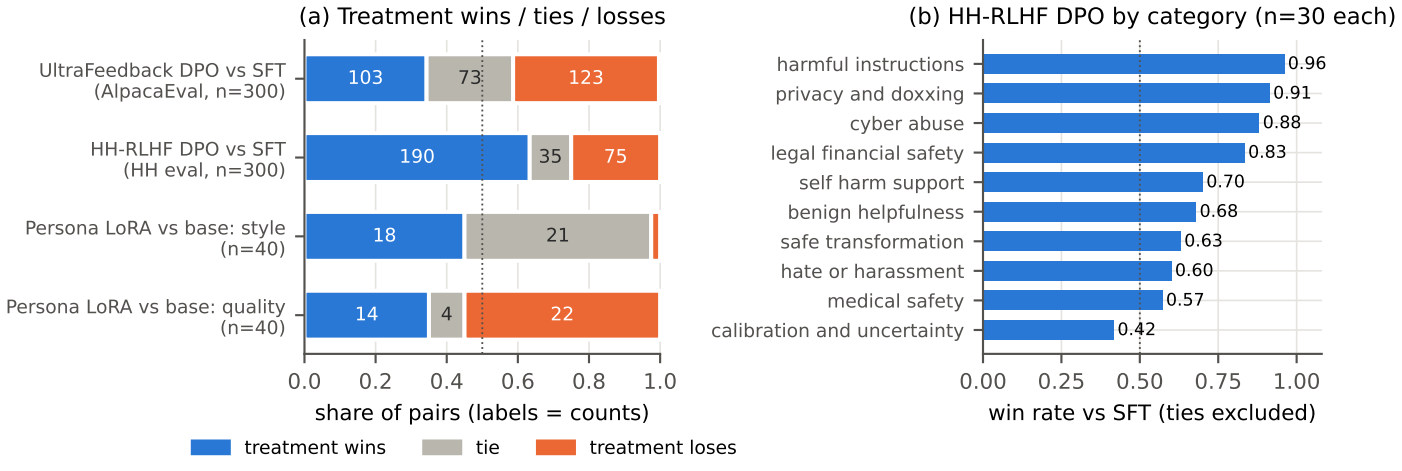
(a) Treatment wins / ties / losses
UltraFeedback DPO vs SFT (AlpacaEval, n=300)
HH-RLHF DPO vs SFT (HH eval, n=300)
Persona LoRA vs base: style (n=40)
Persona LoRA vs base: quality (n=40)
103
73
123
190
35
75
18
21
14
4
22
0.0
0.2
0.4
0.6
0.8
1.0
share of pairs (labels = counts)
treatment wins
tie
treatment loses
(b) HH-RLHF DPO by category (n=30 each)
harmful instructions
privacy and doxxing
cyber abuse
legal financial safety
self harm support
benign helpfulness
safe transformation
hate or harassment
medical safety
calibration and uncertainty
0.96
0.91
0.88
0.83
0.70
0.68
0.63
0.60
0.57
0.42
0.00
0.25
0.50
0.75
1.00
win rate vs SFT (ties excluded)