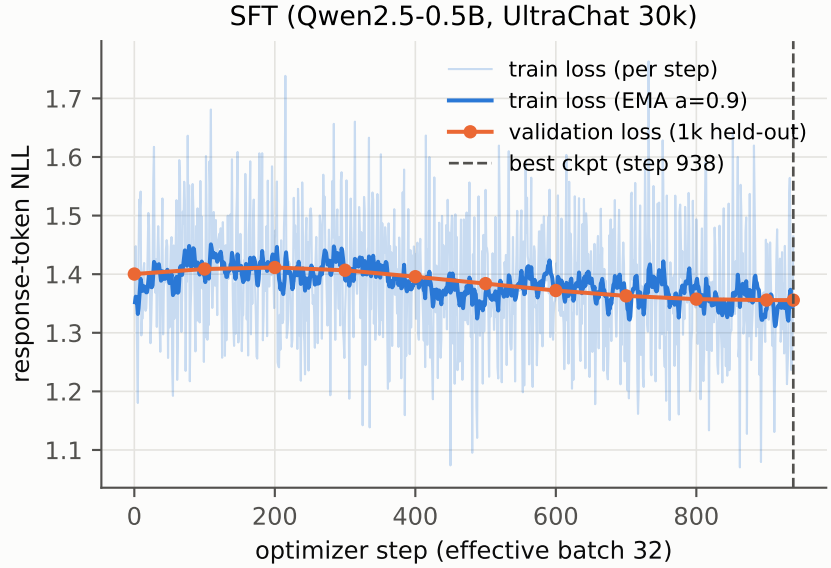
SFT (Qwen2.5-0.5B, UltraChat 30k)
train loss (per step)
train loss (EMA a=0.9)
validation loss (1k held-out)
best ckpt (step 938)
response-token NLL
1.7
1.6
1.5
1.4
1.3
1.2
1.1
0
200
400
600
800
optimizer step (effective batch 32)

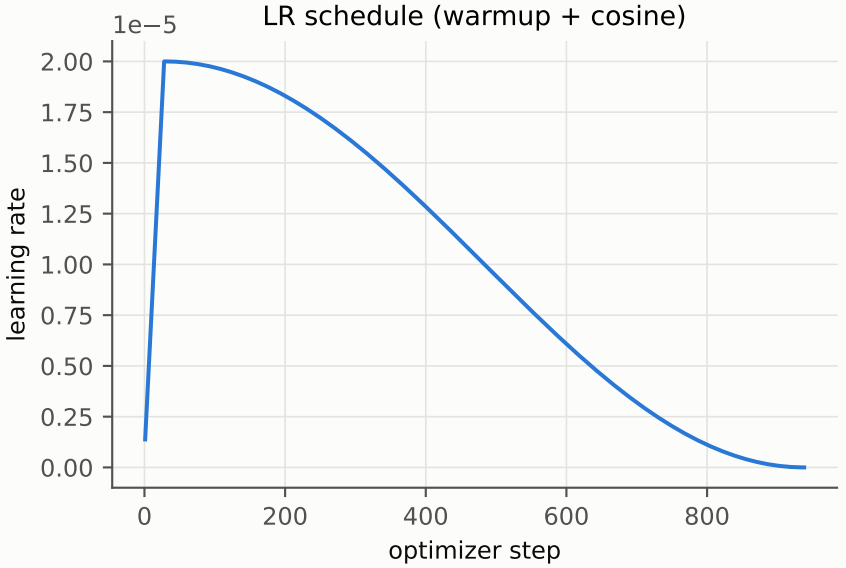
LR schedule (warmup + cosine)
1e−5
learning rate
2.00
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
0
200
400
600
800
optimizer step